

Q1 2026 Earnings Call

Company Participants

- Mark Murphy, Executive Vice President and Chief Financial Officer
- Sanjay Mehrotra, Chairman, President and Chief Executive Officer
- Satya Kumar, Corporate Vice President, Investor Relations and Treasury

Other Participants

- C.J. Muse, Cantor Fitzgerald
- Christopher Danelly, Citi
- Harlan Sur, J.P. Morgan
- Krish Sankar, TD Cowen
- Thomas O'Malley, Barclays
- Timothy Arcuri, UBS
- Vivek Arya, Bank of America Securities

Presentation

Operator

Thank you for standing by, and welcome to Micron's First Quarter 2026 Financial Call. At this time, all participants are in listen-only mode. After the speakers' presentation, there will be a question-and-answer session. (Operator Instructions). As a reminder, today's program is being recorded.

And now I'd like to introduce your host for today's program, Satya Kumar, Investor Relations. Please go ahead, sir.

Satya Kumar {BIO 18241786 <GO>}

Thank you and welcome to Micron Technology's fiscal first quarter 2026 financial conference call. On the call with me today are Sanjay Mehrotra, our Chairman, President and CEO; and Mark Murphy, our CFO.

Today's call is being webcast from our Investor Relations site at investors.micron.com, including audio and slides. In addition, the press release detailing our quarterly results has been posted on the website, along with the prepared remarks for this call.

Today's discussion contains forward-looking statements that are subject to risks and uncertainties. These forward-looking statements include statements regarding our future financial and operating performance, as well as trends and expectations in our business, contractual terms, market, industry, products, and regulatory and other matters. These statements are based on our current assumptions, and we assume no obligation to update these statements.

Please refer to our most recent financial reports on Forms 10-K, Forms 10-Q, and our other filings with SEC for more information on the risks and uncertainties that could cause actual results to differ materially from expectations.

Today's discussion of financial results is presented on a non-GAAP financial basis unless otherwise specified. A reconciliation of GAAP to non-GAAP financial measures can be found on our website.

I'll now turn the call over to Sanjay.

Sanjay Mehrotra {[BIO](#) [1977912](#) <[GO](#)>}

Thank you, Satya. Micron had an outstanding start to fiscal 2026, delivering fiscal Q1 revenue, gross margin, and EPS, well above the high end of our guidance. This financial performance was driven by our strong execution across end markets and products in a tight supply environment.

We achieved a number of records in fiscal Q1. Total company revenue, DRAM and NAND revenue, as well as HBM and data center revenue, and revenue in each of our business units also reached new records. We have completed agreements on price and volume for our entire calendar 2026 HBM supply, including Micron's industry-leading HBM4.

We forecast an HBM TAM CAGR of approximately 40% through calendar 2028, from approximately \$35 billion in 2025 to around \$100 billion in 2028. This \$100 billion HBM TAM milestone is now projected to arrive two years earlier than in our prior outlook. Remarkably, this 2028 HBM TAM projection is larger than the size of the entire DRAM market in calendar 2024. We are excited about our customized HBM4E customer engagements, which offer further differentiation opportunities to us, and we continue to make excellent progress on our HBM roadmap.

Memory is now essential to AI's cognitive functions, fundamentally altering its role from a system component to a strategic asset that dictates product performance from data center to the edge. This structural shift means that system capabilities heavily rely on advanced memory for real-time contextual processing, which is vital for achieving autonomous and intelligent behaviors in AI data centers, as well as in applications ranging from self-driving cars to advanced medical diagnostics.

With our technology leadership, differentiated product portfolio, strong operational execution, and solid balance sheet, Micron is in the best competitive position in its history and is one of the semiconductor industry's biggest enablers of AI.

We anticipate substantial new records in revenue, gross margin, EPS, and free cash flow for both the second quarter and the full fiscal year 2026, and we expect our business performance to continue to strengthen through the year.

Sustained and strong industry demand, along with supply constraints, are contributing to tight market conditions, and we expect these conditions to persist beyond calendar 2026. We are making progress with customers in our discussions for multi-year contracts with specific commitments. Simultaneously, we are focused on maximizing our production output from our current footprint, ramping our industry-leading technology nodes, and investing in new cleanroom space to add to our supply capability.

Micron's technology leadership is foundational to our strong competitive position. Micron has led the industry for four consecutive technology nodes in DRAM and three nodes in NAND, with progressively faster yield ramps in every node. Our 1-gamma DRAM node is ramping well. 1-gamma will be the primary driver of our DRAM bit growth in calendar 2026 and will be the majority of our bit output in the second half of the calendar year.

Looking beyond 1-gamma, development is underway for 1-delta and 1-epsilon nodes, which will feature innovations that we expect to extend our differentiation and technology leadership. In NAND, we are ramping our G9 node with robust yield ramps across both data center and client SSDs. Our QLC NAND mix, including G9 QLC, reached a record high during the quarter.

Technology transitions to G9 will be the primary driver of our NAND bit growth in calendar 2026, and we expect it to become our largest NAND node later in fiscal 2026.

I'm pleased to report that calendar 2025 is a record year for Micron in terms of both internal and customer quality measures, positioning us well to deliver for our customers as the memory industry's quality leader. As our products are increasingly integrated into higher value applications, our leadership in quality is becoming a more important differentiator.

Turning to our end markets. As the world's leading technology companies advance toward artificial general intelligence and transform the global economy, our customers are committing to an extraordinary multi-year data center build out. This growth in AI data center capacity is driving a significant increase in demand for high performance and high capacity memory and storage.

Server unit demand has strengthened significantly, and we now expect calendar 2025 server unit growth in the high-teens percentage range, higher than our last earnings call outlook of 10%. We expect server demand strength to continue in 2026. Server memory and storage content and performance requirements continue to increase generation-to-generation. Micron has a differentiated portfolio of high-value data center solutions to address these requirements, including our HBM, high-capacity server memory solutions, and data center SSDs.

Micron's HBM4, with industry-leading speed over 11 gigabits per second, is on track to ramp with high yields in the second calendar quarter of 2026, consistent with our customers' product ramp plans. Our HBM4 uses advanced CMOS and advanced metallization process technologies on the base logic die and DRAM core dies, which are designed and manufactured in-house. This, along with our unique HBM design, packaging, and test capability, enables Micron's industry-leading performance and low-power leadership.

Micron pioneered the adoption of LPDRAM in the data center. Micron's low-power DRAM server modules consume one-third the power of DDR DRAM server modules. Building on this leadership, we have sampled our 192 gigabyte LP SOCAMM2 product, which enables a 50% increase in capacity per module and a rack-scale LPDRAM density of over 50 terabytes.

Our data center NAND portfolio revenue exceeded \$1 billion in fiscal Q1, and we are seeing strong momentum across our data center SSD portfolio, enabled by our leadership NAND technology. In the performance SSD category, Micron has introduced the world's first PCIe Gen6 SSD, leveraging our G9 NAND. We are seeing rapidly increasing qualification commitments for this product, including at hyperscalers.

In mainstream storage, our SSDs based on G9 NAND are already seeing robust demand in the first quarter of calendar 2026. In capacity storage, our QLC-based 122 and 245 terabytes G9 SSDs are entering qualification at multiple hyperscale customers.

PC demand continues to be driven by Windows 10 end-of-life and AI PCs. We forecast PC unit sales to grow high-single digit percentage range in calendar 2025, above our prior expectations, provided in our last earnings call of mid-single digits. As we look ahead into 2026, we expect these demand drivers to continue while memory supply constraints may affect some PC unit shipments. Micron has completed multiple OEM qualifications of our 16-gigabit, 1-gamma-based DDR5 and our G9-based PCIe Gen4 QLC SSDs.

Turning to mobile. Smartphone unit volumes in calendar 2025 are on track to grow in the low-single digit percentage range. AI is driving memory content growth. The shipment mix of flagship smartphones with 12 gigabyte of DRAM increased to 59% in calendar Q3, more than twice the level from a year ago. Micron is accelerating innovation across our mobile DRAM portfolio.

In fiscal Q1, we began sampling our breakthrough 1-gamma 16-gigabit LPDDR6 product to leading OEM and ecosystem partners, marking a major milestone in next-generation memory technology. LPDDR6 will power AI at the edge, delivering over 50% higher performance and improved power efficiency for flagship smartphones and AI PCs. Micron also sampled our 1-gamma LP5X 24-gigabit product and began volume shipments of the previously announced 1-gamma LP5X 16-gigabit product to multiple OEMs.

Turning to auto, industrial and embedded. In automotive, L2 plus and L3 adoption is driving robust demands today, and our customers' roadmaps indicate a significantly higher memory content in fully automated vehicles. Micron is uniquely positioned for growth with our differentiated product portfolio and automotive market share leadership.

Our ASIL-rated LPDDR5X and UFS 4.1 NAND products, optimized for automotive and advanced robotics that include bandwidth-enhancing features, are seeing strong demand and have already secured billions of dollars in design wins.

In industrial, demand continues to strengthen, driven by the growing adoption of autonomous systems across various applications. Long-term demand trend trajectory remains robust for memory and storage in industrial applications, such as in factory automation, aerospace and defense, humanoid robotics, edge networking, and video surveillance. Across both auto and industrial markets, LPDDR4X and DDR4 are also experiencing strong demand, and we are making investments to provide long-term supply from our Manassas, Virginia fab.

Now turning to our market outlook. Over the last few months, our customers' AI data center build-out plans have driven a sharp increase in demand forecasts for memory and storage. We believe that the aggregate industry supply will remain substantially short of the demand for the foreseeable future. The dramatic increase in HBM demand is further challenging the supply environment due to the 3-to-1 trade ratio with DDR5, and this trade ratio only increases with future generations of HBM. Additional cleanroom space is necessary to address this increased demand and lead times for cleanroom build-out or lengthening across geographies.

Together, these demand and supply factors are driving tight industry conditions across DRAM and NAND, and we expect tightness to persist through and beyond calendar 2026. Calendar 2025 DRAM and NAND industry bit demand growth expectations are higher than in our last earnings call outlook. We now expect calendar 2025 DRAM bit demand growth to be in the low 20% range versus high-teens previously.

We expect 2025 NAND bit demand growth to be in the high-teens percentage range versus low-to mid-teens previously. We expect calendar 2026 industry DRAM and NAND bit shipment growth to be constrained by industry supply. We expect both DRAM and NAND calendar 2026 industry bit shipments to increase around 20% from 2025 levels.

Micron is working hard to support our customers' demand during this time, and we expect to grow our DRAM and NAND bit shipments approximately 20% in calendar 2026. Despite significant efforts, we are disappointed to be unable to meet demand from our customers across all market segments.

Micron plans to increase our fiscal 2026 CapEx to approximately \$20 billion versus our prior estimate of \$18 billion. This increase will primarily support our HBM supply capability and also our 1-gamma supply in calendar 2026. We are pulling in equipment orders and accelerating installation timelines to maximize output capability.

Micron is also investing across our global manufacturing footprint to add supply to support longer-term demand. We are seeing an enthusiastic customer response to our planned U.S. supply. We are pulling in our first Idaho fab timeline, and we now expect first wafer output in mid-calendar 2027, earlier than our prior expectation of second half calendar 2027.

Earlier this year, we announced our plans for the second Idaho fab, which will begin construction in 2026 and be operational by the end of 2028. We are making good progress on securing necessary permits for our New York site and appreciate the partnership with the state of New York and the Trump administration. We plan to break ground on our first New York fab in early calendar 2026, which we expect will provide supply in 2030 and beyond.

In Japan, with the support of METI, we are making technology and manufacturing investments. We are enabling future DRAM technology transition in coordination with our Boise R&D team. We are also adding cleanroom space in our Hiroshima fab to support these advanced nodes, which will increase production scale and optimize fab economics.

In Singapore, our HBM advanced packaging facility is on track to contribute meaningfully to our HBM supply in calendar 2027. As HBM becomes a part of our Singapore manufacturing footprint, we expect opportunities for synergies between NAND and DRAM production. We are pleased with the progress on our assembly and test facility in India, which has initiated pilot production and will ramp in 2026.

As we make progress on our strategic manufacturing initiatives, we will continue to be responsive to the market environment and disciplined with our CapEx plans.

I will now turn it over to Mark for our fiscal Q1 financial results and outlook.

Mark Murphy {[BIO 16136048](#) <GO>}

Thank you, Sanjay, and good afternoon, everyone. Micron delivered strong financial results for the fiscal first quarter, with revenue, gross margin, and EPS all exceeding the high end of our guidance.

During the quarter, we generated record free cash flow, reduced our debt, and returned to net cash. Total fiscal Q1 revenue was \$13.6 billion, up 21% sequentially and up 57% year-over-year, setting a quarterly record for the third consecutive quarter. We saw sequential revenue growth across all our business units.

Fiscal Q1 DRAM revenue was a record \$10.8 billion, up 69% year-over-year, and represented 79% of total revenue. Sequentially, DRAM revenue increased 20%. Bit shipments were up slightly and prices increased approximately 20%, driven by tight industry DRAM supply, pricing execution, and favorable mix.

Fiscal Q1 NAND revenue was a record \$2.7 billion, up 22% year-over-year, and represented 20% of Micron's total revenue. Sequentially, NAND revenue increased 22%. NAND bit shipments increased in the mid-to-high single digit percentage range, and prices increased in the mid-teens percentage range, driven by tight NAND industry supply, pricing execution and favorable mix.

The consolidated gross margin for fiscal Q1 was 56.8%, up 11 percentage points sequentially. This improvement was driven by higher pricing, with strong cost execution and favorable mix.

Now turning to quarterly financial performance by business unit. Cloud Memory Business Unit revenue was a record \$5.3 billion, and represented 39% of total company revenue. CMBU revenue was up 16% sequentially, driven by an increase in bit shipments and higher prices. CMBU

gross margins were 66%, higher by 620 basis points sequentially, supported by cost execution and higher pricing.

Core Data Center Business Unit revenue was a record \$2.4 billion, and represented 17% of total company revenue. CDBU revenue was up 51% sequentially, driven by robust bit shipments and higher pricing. CDBU gross margins were 51%, up 990 basis points sequentially, supported by higher pricing and cost execution.

Mobile and Client Business Unit revenue was a record \$4.3 billion, and represented 31% of total company revenue. MCBU revenue was up 13% sequentially, driven by higher pricing, partially offset by lower bit shipments. MCBU gross margins were 54%, up 17 percentage points sequentially, driven primarily by higher pricing.

Automotive and Embedded Business Unit revenue is a record \$1.7 billion and represented 13% of total company revenue. AEBU revenue was up 20% sequentially, driven by higher bit shipments and higher pricing. AEBU gross margins were 45%, up 14 percentage points sequentially, driven primarily by higher pricing.

Operating expenses in fiscal Q1 were \$1.3 billion, up \$120 million quarter-over-quarter and in line with our guidance range. The sequential increase was driven by higher R&D expenses in support of technology and product development on our new DRAM and NAND technology nodes.

We generated operating income of \$6.4 billion in fiscal Q1, resulting in an operating margin of 47%, up 12 percentage points sequentially and 20 percentage points year-over-year. Fiscal Q1 taxes were \$977 million on an effective tax rate of 15.1%. Non-GAAP diluted earnings per share in fiscal Q1 was \$4.78, with 58% sequential growth and 167% versus the year-ago quarter.

Turning to cash flow and capital expenditures. In fiscal Q1, operating cash flows were \$8.4 billion and capital expenditures were \$4.5 billion, resulting in free cash flow of \$3.9 billion. Fiscal Q1 free cash flow was a quarterly record, exceeding our prior record in fiscal Q4 2018 by over 20%. Ending inventory for fiscal Q1 was \$8.2 billion, down \$150 million sequentially, with days of inventory at 126. DRAM inventory days remain tight and below 120 days.

On the balance sheet, we held \$12 billion of cash and investments at quarter end, and maintained \$15.5 billion of liquidity when including our untapped credit facility. In fiscal Q1, we repurchased \$300 million of shares as permitted by the terms of the CHIPS agreement.

During the quarter, we also reduced debt by \$2.7 billion, paying off \$1 billion balance of term loans and redeeming \$1.7 billion of senior notes. We closed the quarter with \$11.8 billion of debt and a net cash balance over \$250 million. Through the fiscal year, we expect to further strengthen our balance sheet as we generate additional free cash flow.

Before turning to our outlook, I would like to share an update on how we are benefiting from AI use across Micron. Today, over 80% of our professional workforce actively uses GenAI, with total usage up tenfold since last year. In manufacturing, integrating AI into yield and quality management has cut root cause identification time by half in cases. Our coding teams are realizing productivity gains of 30% or more using agentic AI.

In R&D, GenAI is accelerating development by reducing cycle times and design verification, product validation, issue triage, and root cause analysis. Across business functions, GenAI is broadening automation opportunities, and we are deploying conversational analytics to accelerate and improve decision-making. We expect Micron's use of AI across the enterprise to further strengthen our competitiveness in the coming years.

Now turning to our outlook for the fiscal second quarter. Industry demand is greater than supply for both DRAM and NAND. We expect higher price, lower cost, and favorable mix to all contribute to gross margin expansion in Q2. Operating expenses for fiscal Q2 are projected to be approximately \$1.38 billion. As mentioned last quarter, Micron's fiscal Q4 2026 OpEx will also reflect the effect of an additional work week in this 53-week fiscal year. We expect a fiscal Q2 and fiscal year 2026 tax rate of around 15.5%.

Micron is investing in a disciplined manner across our global manufacturing footprint to better meet demand. To address tight supply-demand conditions extending beyond 2026, we now project our capital spending in fiscal 2026 to be approximately \$20 billion, weighted to the second half of the fiscal year. We expect free cash flow to strengthen in fiscal Q2, and we expect to generate significantly higher free cash flow year-over-year in fiscal 2026.

Any impacts that may occur due to potential new tariffs are not included in our guidance. With all these factors in mind, our non-GAAP guidance for fiscal Q2 is as follows. We expect revenue to be a record \$18.7 billion, plus or minus \$400 million; gross margin to be in the range of 68%, plus or minus 100 basis points; and operating expenses to be approximately \$1.38 billion, plus or minus \$20 million. Based on a share count of approximately 1.15 billion shares, we expect EPS to be a record \$8.42 per share, plus or minus \$0.20.

I'll now turn it over to Sanjay to close.

Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

Thank you, Mark. AI-driven demand is here, and it is accelerating. And Micron is capturing these opportunities with the best competitive position in its history. This success is built on the strength of our global team, and I want to thank our team members worldwide for their hard work and dedication. We are in the most exciting time in Micron's history, and the best is yet to come.

We will now open for questions.

Satya Kumar {[BIO 18241786 <GO>](#)}

Operator, can you queue up the questions?

Questions And Answers

Operator

(Question And Answer)

Yes, sir. And I show our first question comes from the line of Timothy Arcuri from UBS. Timothy Arcuri from UBS, your line is open.

Q - Timothy Arcuri {[BIO 3824613 <GO>](#)}

Thanks a lot. Sanjay, I wanted to ask you about customer LTAs. I know we're hearing about D5 that's being bundled with HBM, and in some cases even NAND. So can you just talk about these LTAs? I know it sounds like these are stretching out through '26, and in some cases even into '27. I've even heard of some stuff into '28. So can you talk about the nature of these LTAs? And then I had a follow-up as well. Thanks.

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

These are multi-year contracts that we are in discussions with several of our key customers. And these contracts, of course, involve DRAM as well as NAND. And with respect to terms, of course, these contracts that we are under discussion for are very different from prior LTAs. They have specific commitments in them and much stronger contract structure. And beyond that, I can't be giving you specifics at this point. Of course, in the future, if and when appropriate, we'll be sharing further details.

Q - Timothy Arcuri {BIO 3824613 <GO>}

Thanks. And then, Mark, I wanted to ask you about CapEx. So, you took it up to \$20 billion net, but it still seems, I mean, you're not guiding all fiscal '26 revenues. So, we don't really know what the capital intensity number is. But it seems like it's like 25% to 30%, which is a little below your 35% metric that you usually think of. So, is that because you're constrained because of fab space? And can you just talk about, does that sort of like roll into fiscal '27, where we would see CapEx up more near that 35% range? Thanks.

(Technical Difficulty)

Operator

Thank you. And I show our next question in the queue comes from the line of C.J. Muse from Cantor Fitzgerald. Please go ahead.

Q - C.J. Muse {BIO 6507553 <GO>}

Satya, back to you.

Operator

C.J. Muse, your line is open.

Q - C.J. Muse {BIO 6507553 <GO>}

Yes. They didn't answer the prior question. Operator?

Operator

Yes, sir.

Q - C.J. Muse {BIO 6507553 <GO>}

Can you turn it back to management so they can answer the prior question?

Operator

Sure. Management's line is open. Please proceed. I believe the speaker's lines are muted at this time. Ladies and gentlemen, please continue to stand by. Your call will resume momentarily.

A - Mark Murphy {BIO 16136048 <GO>}

Operator, can you hear us?

Operator

Yes, we can now. You're unmuted, sir. Please proceed.

A - Mark Murphy {BIO 16136048 <GO>}

Okay. I'm not sure. What happened there? We were on mute here. So, I just want to make sure that you heard Sanjay's response. Correct?

Q - Timothy Arcuri {BIO 3824613 <GO>}

We did not.

A - Mark Murphy {BIO 16136048 <GO>}

You did not hear Sanjay's response. Okay. Thanks, Tim. So, Sanjay?

Q - Timothy Arcuri {BIO 3824613 <GO>}

We heard -- I heard Sanjay's response to the first question, but I didn't hear your response, Mark, to my question on CapEx and on capital intensity.

A - Mark Murphy {BIO 16136048 <GO>}

Okay, Tim. Thank you. So you're right, Tim, we're not providing a full year revenue guide. We did indicate that our CapEx was going up in calendar -- in fiscal '26. And substantial part of that CapEx is for DRAM and specifically HBM and the 1-gamma and 1-gamma ramp.

I would say that from '25 to '26, the plan is roughly to double the brick-and-mortar construction CapEx. And at this time, we would expect '27 CapEx to be up and -- but I want to emphasize that Micron is going to remain disciplined on CapEx growth to support bit demand and that -- or, bit supply, and that supply will be in line with demand.

As your question of capital intensity, our capital intensity, of course, is dropping as the market conditions remain very constructive, and of course, we are working to be very efficient with our capital spend.

Q - Timothy Arcuri {BIO 3824613 <GO>}

Okay. Mark, thank you.

Operator

Thank you. And I show our next question comes from the line of C.J. Muse from Cantor Fitzgerald. Please go ahead.

Q - C.J. Muse {BIO 6507553 <GO>}

Yes. Good afternoon. Thank you for taking the question. I guess, Mark, to follow-up on the prior question around CapEx, and the relative growth seems very conservative in the backdrop that we're in, and it feels like that you're just sitting here without cleanroom space, and it also doesn't sound like you're meaningfully pulling in cleanroom. So, can you talk about the philosophy there, and I guess what I'm taking away from your commentary is that you're being very conservative and judicious with adding capacity here.

A - Mark Murphy {BIO 16136048 <GO>}

Well, I would say that we've been indicating issues with supply for several quarters that we were working inventories down, that node transitions were going to be the principal source of supply

growth in fiscal '26, and that's exactly what's happening, and we know that cleanroom space takes time, and the HBM growth, which is only picked up with AI-driven demand, has further pressured supply.

So, there's no near-term solution, as we said in the prepared remarks, the entire industry we expect to be short to demand, and we're no different in that case. But we are moving quickly to do our best to provide customer supply. And we have pulled in tools. We have accelerated construction in Idaho. We are doing everything we can within our existing footprint and near-term capacity expansions to deliver supply. And so we provided a bit growth number for '26, and that is supply constrained.

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

And I'll just add that, of course, we are continuing to make the investments in technology transitions in our existing footprint. And as we highlighted that 1-gamma node will be a majority driver of our supply growth in '26. And of course, you have seen us make investments not only in technology transitions, but also in greenfield capacity, but also enabling greater technology production capability in our existing cleanrooms in Japan. And we are making the necessary investments there as well.

So, of course, we remain disciplined, but we are very much focused and trying to work hard toward increasing our supply there. And pleased with our plans for Idaho 1, Idaho 2, and of course, New York as well. And of course, in the short term, very much focused on maximizing production efficiencies, maximizing our production output from the existing footprint as well.

But yes, I mean, demand fundamentals are pretty strong, driven by AI from data center to edge with the build out of our customers. And supply is significantly short. And I would say that in the medium term, we are only able to meet about 50% to two-thirds of our demand from several key customers. So we remain extremely focused on trying to increase the supply here and making the necessary investments.

Q - C.J. Muse {[BIO 6507553](#) <GO>}

Very helpful. And then I guess as a follow-up for gross margins, obviously, the guide it's quite stellar, but curious as you go through calendar '26, how should we think about cost down across both DRAM and NAND? And as you transition from 3E to 4, is there anything that we should keep in mind or we should think about and contemplate in our models where there might be higher costs temporarily giving yields or whatnot? Thanks so much.

A - Mark Murphy {[BIO 16136048](#) <GO>}

Yes, so C.J., our cost execution has been very good across both DRAM and NAND. Of course, we're getting some volume leverage, but spend control has been very good. Yields have been good. We do have some start-up costs coming in for the new fabs, new construction across the network. That starts to come in second half of '26 and into '27. But at these size of the business at these levels, it's a relatively small impact on margin.

We're not going to provide cost guidance for the rest of the year as it depends on many factors, including mix. But to our earlier question, I can say that we've talked about ramping 1-gamma DRAM and G9 NAND at length for supply in '26. And those ramps are proceeding well and will be a tailwind or a cost as these nodes ramp.

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

And regarding your question on HBM3E and HBM4, as we have said, we'll be beginning to ramp production of HBM4 in Q2 timeframe in line with our customer demands. And of course, our HBM4 is progressing extremely well, very pleased with our product, industry-leading product with the highest performance of over 11 gigabits per second.

And so, I mean, that's the highest performance and very pleased with its overall yield ramp. And we expect our HBM4 to be expecting, having a faster yield ramp than our HBM3E. And of course, our mix of HBM3E and HBM4 during '26 will be very much based on our overall customer demands. And we will have both of these products with strong profile in our calendar '26 revenue.

Q - C.J. Muse {[BIO 6507553 <GO>](#)}

Thank you.

Operator

And I show our next question in the queue comes from the line of Harlan Sur from JPMorgan. Please go ahead.

Q - Harlan Sur {[BIO 6539622 <GO>](#)}

Good afternoon and great job on the quarterly execution. Just over the past three to four months as we track the different ASIC AI XPU programs, there's been a significant upper revision on ASIC XPU volume shipments next year. You have Google TPU, AWS Trainium, and so on, right? And all of these XPUs are still going to be using HBM3E.

Have you -- has the team seen this near-term positive dynamic in your order book for 3E? And given that you're fully contracted for calendar '26, like what -- and what appears to be growing upside to next year's view, how is the team -- how is the Micron team going to try and manage this upside dynamic in 3E alongside a strong HBM4 demand profile?

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

So as I mentioned, Harlan, earlier that, of course, 2026 will have a mix of HBM3E and HBM4. And we have shared with you in the past that we are engaged with multiple customers with the entire ecosystem here of HBM customers and very much engaged with them. And they will all contribute to our strong year-over-year growth in revenue in '26. And of course, that will be made up of both HBM3E and HBM4.

So as I said, I mean, we will continue to manage the mix of the two based on customer requirements. I can tell you that in 2026, supply on HBM will be tight. Non-HBM DRAM will be tight as well. So we are continuing to see, as we have highlighted in our prepared remarks, tightening supply environment. And of course, we see strong year-over-year growth in '26 for our HBM. We today upped our revenue forecast for HBM.

We highlighted that by 2028, we expect it to be \$100 billion TAM. And that's two years ahead of our prior outlook. So, of course, HBM is on a good trajectory. And what I can also tell you is that customers, as their architectures are evolving, as their platforms are evolving, and these are customers across the ecosystem, of course, they are requiring more and more HBM. I mean, the value of memory in terms of ability to deliver the AI capabilities and the functionality and the performance, memory is critical. And more HBM is required, and that's across the various AI platforms in the industry.

Q - Harlan Sur {BIO 6539622 <GO>}

I appreciate that, Sanjay. And then after two to three quarters of enterprise SSD, sort of muted trends, I believe the team saw a strong acceleration in the business, right, according to some of the third-party research estimates. I think your enterprise SSD business grew like 25% sequentially in the most recent quarter, right? So you're the number three market share leader amongst eight or nine competitors, right? Very strong share position.

Given, I would assume, increasing demand trends here, expanding lead times, is the Micron team also entering into long-term supply agreements with your eSSD customers? And then secondarily, I mean, SSD demand, is it more tied to expansion in inferencing workloads as customers aggressively move to monetization? In other words, is storage intensity higher on inferencing versus training workloads?

A - Sanjay Mehrotra {BIO 1977912 <GO>}

So with respect to enterprise SSDs, really very proud of our engineering and business teams, and of course, our sales teams, in terms of our customer engagement and the strong momentum and the share gains that we have with our enterprise SSD. Of course, our enterprise SSDs are a big part -- an important part, let me say, of our data center strengthening mix. Of course, DRAM is continuing to increase in mix to our data center, but data center SSDs are an important part of our overall revenue mix.

And we expect to continue to focus on share gains with our strong SSD roadmap, customer engagements, and great quality that we provide to the customers. And as I mentioned, that our multi-year contract that we are in discussions with, with our several key customers, our SSDs -- our data center SSDs, are also part of that. And let me tell you that these multi-year contracts are not just about data center customers. They also are about multiple customers across our market segments here.

And in terms of your questions on is the SSD requirement growing with inferencing versus training, what I can tell you is that in the data center AI applications, I mean, as the generative AI moves to more and more video, of course that drives greater demand for more SSDs as well. So, I mean the rapid evolution of AI from training to inferencing and the rapid evolution of AI models and applications, they're all driving greater growth of enterprise SSDs, yes, fueled by GenAI.

Q - Harlan Sur {BIO 6539622 <GO>}

Thanks, Sanjay.

Operator

And I show our next question comes from the line of Tom O'Malley from Barclays. Please go ahead.

Q - Thomas O'Malley

Hey, guys. Thanks for taking my question. Sanjay, you've been helpful in the past about kind of talking about Micron's ramp in HBM and then also the ramp of the total market. You gave some new color on HBM with the \$35 billion moving at a 40% CAGR, but Micron specifically, I was curious if you could give us any color on percentage of the DRAM business today that's HBM from a dollar's perspective.

And then on share, as you move into next year, obviously there's a large competitor that is looking to become more competitive at 3E. We haven't heard anything on 4 yet. How do you feel about your competitive positioning into next year? And do you think that you're going to make any strategic decisions differently based on the public certification of their memory in the next couple of months? Thank you very much.

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

We feel very good about our competitive position. We feel very, very good about our product and our HBM4 product that we have highlighted as industry leading performance over 11 gigabit per second, the best specifications in the industry with our performance. And of course, we feel very good about the power consumption in our products as well.

In the past, we have shared with you that our HBM3E is 30% lower power than any of the competitors in the industry, and we are maintaining that momentum of low power, which you know is, in data center applications, is very important. So, we are maintaining our performance and power, and of course, our strong capacity position with HBM4 roadmap as well.

So, we feel very good about our competitive position, about our roadmap, and our roadmap going beyond HBM4 for the future years, beyond '26 as well. And we are very proud of our team's ability to execute successfully over the course of last several quarters in terms of ramping up production capabilities of HBM3E. And we shared with you that in CQ3, we reached our share of HBM3E to be in line with our -- I mean, our HBM share to be in line with our DRAM share.

And we have always highlighted that as we have reached that share, we will -- particularly in this tight supply environment, we'll be managing the mix of our HBM as well as our non-HBM. All of it is in high demand, and HBM as well as non-HBM has strong profitability. So, looking at our strategic customer relationships as well as our overall profitability goals and growth objectives, we'll continue to manage that mix between HBM and non-HBM.

But of course, HBM is growing and we have highlighted that how we expect the TAM to be \$100 billion by 2028, a couple of years ahead of our prior projection. And of course, we will grow our HBM as well. 2026 will see a strong year-over-year growth in our HBM. And in this very tight supply environment, as I mentioned, that the gap between the demand and supply for all of DRAM, including HBM, is really highest that we have ever seen, and I quantified it earlier as well.

So, in this environment, of course, working closely with our customers, we are continuing to manage our mix of the product here.

Q - Thomas O'Malley

Thank you, Sanjay. And just --

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

As we focused on also continuing to increase supply to better address our customers' demand requirements.

Q - Thomas O'Malley

Perfect. And then just as a follow-up, you've said historically kind of the \$8 billion run rate. If you look at November and February, you're taking up the total TAM, but any color specifically on HBM contribution in the November quarter and what you're expecting in the guide?

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

We are not really providing those specifics here in terms of the breakout. I mean, we highlighted that in FQ1, our HBM revenue was a record. We did highlight that. And beyond that, we are really not going to be providing the specifics in terms of revenue. And just would tell you again that year-over-year in '26, we'll be seeing a strong growth in our HBM revenue.

Q - Thomas O'Malley

Thank you, Sanjay.

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

And again, our product is very well positioned. So, I mean, that's a very good place to be in, in terms of managing overall mix of the business.

Operator

Thank you. And I show our next question comes from the line of Krish Sankar from TD Cowen. Please go ahead.

Q - Krish Sankar {[BIO 16151788 <GO>](#)}

Yes. Hi. Thanks for taking my question and congrats on the phenomenal results and guidance. My first question is for Mark. I know you spoke about the sustainability for the next year. I'm kind of curious how to think about gross margins beyond the February quarter, like into May. Is it going to improve or sustain at these levels? How to think about the gross margins? I had a follow-up for Sanjay.

A - Mark Murphy {[BIO 16136048 <GO>](#)}

Thanks, Krish. We're not guiding margins beyond Q2. We did guide a record Q2, as you know 68%, 11 points sequential improvement, 7 points better than the previous record.

We did indicate that we would -- our business would strengthen through the year. And so we do believe margins can be up. We do believe that they will be up for DRAM and NAND. Now keep in mind that at these high gross margin levels mathematically we get less in gross margin percent for the same increase in price. So yes, we would expect gross margins to expand beyond fiscal Q2, but we would expect that growth to be more gradual than what we've seen in the last couple of quarters -- or, the first quarter and the second quarter guide.

Q - Krish Sankar {[BIO 16151788 <GO>](#)}

Got it. Thanks a lot, Mark, that's super helpful. And then Sanjay --

A - Mark Murphy {[BIO 16136048 <GO>](#)}

And Krish, just one last thing. We have indicated that strengthen through the year because we believe this constructive market environment will remain so through the year, these favorable market conditions. But we also -- we're executing very well on cost. And as Sanjay mentioned earlier, we're deploying the bits to the valuable part of the market and where we can serve our customers best.

Q - Krish Sankar {[BIO 16151788 <GO>](#)}

Got it. Super helpful, Mark. Thanks for that. And Sanjay, just as a follow-up to the earlier question, I understand you don't want to like put some boundary conditions, but a year ago, you totally nailed it when you said you're going to get like HBM market share close to your DRAM market share, DDR market share. So when you talk about, like 2028 CAGR of 40%, \$100 billion TAM, how to think about Micron's HBM market share in that realm? Should we assume it's going to be in the low-20s percent, so that spectrum, or it's going to be lower? Thank you.

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

So, Krish, again, we are not really going to be specifying the share. As we have said, we will be managing the mix of the business between HBM as well as our non-HBM. It's like any other product in our portfolio that, when you have a strong product roadmap across the portfolio, then, of course, we manage the mix across our portfolio with all the strategic reasons and customer relationships in mind.

So, Krish, we are not really going to break that down. And all I would say is that in this current industry environment which we see as durable industry fundamentals in the foreseeable future, we are in a very good position with all the tailwinds of our product portfolio, and, of course, the increasing value of memory across the board, of course, HBM, but also non-HBM in data center and other markets, we will just remain very focused on managing the mix of our business and of course, managing for the best in the mid-term as well as keeping in mind the longer-term.

Q - Krish Sankar {[BIO 16151788 <GO>](#)}

Thanks a lot, Sanjay. I appreciate it.

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

Thank you.

Operator

One moment for our next question. And our next question comes from the line of Chris Danely from Citi. Please go ahead.

Q - Christopher Danely {[BIO 3509857 <GO>](#)}

Hey, thanks, guys. So, I just wanted to dig in on these long-term customer contracts you guys are negotiating. Can you give us any more sense of when you think you'll be able to sign these and then maybe just talk about what the holdup is? Is it just the unprecedented length or size? And given that the AI companies are asking for so much of your capacity, are you able to get them to more or less contribute to the building of a new fab? Thanks.

A - Sanjay Mehrotra {[BIO 1977912 <GO>](#)}

So, we will not really get into the specifics around our contract discussions with our customers, but again, I will highlight a couple of important factors that customers are concerned about long-term access to adequate memory in the environment that we are heading into. And that's leading to constructive dialogues with several key customers and across our multiple markets in terms of their supply as well as other important specific commitments related to these longer-term estimates.

So not getting into the specifics, but as I highlighted, our contract structures that we are discussing are not like anything before. They are far stronger contract structures with specific

commitments. And of course, different from prior contracts is that those used to be like one-year contracts, and these are multi-year in nature as well. And as we look at addressing the customer discussions, of course, we have to look at our overall supply. And I have mentioned to you that in the mid term -- medium term, we are only able to meet half to two-thirds of the demand from our several key customers.

So all of that as we are managing our customer relationships and keeping our strategic objectives in mind, all of that has to be taken into account as we manage our contract discussions. But really, not getting into the specifics here.

Q - Christopher Danely {[BIO 3509857](#) <GO>}

That's still very helpful, Sanjay. Thanks. And for my follow-up, just a question on HBM and the pricing there. So given that the demand is so strong, I think you said you're sold out for '26. Are you guys locked into a set price, or can you let that price more or less float a little bit given how strong demand is like DDR5 does, for example?

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

We are really pleased with our product position and our ability to work with our customers as we highlighted in our prepared remarks that our HBM for 2026 is sold out in terms of volume and our negotiations with customers have been completed for calendar year 2026 for volume as well as pricing.

And as we have always highlighted that our HBM has strong profitability and of course, very much focused on ROI and our non-HBM business also clearly has healthy profitability as reflected in the results that we produced as well as in the guidance that we have provided here.

Q - Christopher Danely {[BIO 3509857](#) <GO>}

Great. Thanks, Sanjay. And congrats again on the results.

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

Thank you.

Operator

Thank you. And I show our next question comes from the line of Vivek Arya from Bank of America Securities. Please go ahead.

Q - Vivek Arya {[BIO 6781604](#) <GO>}

Thanks for taking my questions. Sanjay, I'm curious, at what point does increasing memory price impact demand for electronics? If you set aside the data center and the AI market, do you see some elasticity? Do you see any impact on demand as you look into 2026 for more consumer and kind of traditional enterprise products? How does that shape where memory pricing can go next year?

A - Sanjay Mehrotra {[BIO 1977912](#) <GO>}

We have highlighted in our prepared remarks that, in some of the consumer markets, I mean, some of the unit demand may get impacted, given semiconductor prices here, given memory prices here. And of course, some of the customers may have, for example, in smartphone and

PCs, they may have some mix adjustments in their portfolio as well to address available supply to them.

But these are accounted for in our forecast that we have. And so, I mean, some of the possible impact on unit demand and some of the customer mix changes have been accounted for in our forecast. And we, of course -- even then, we see a very, very tight supply environment here and the large gap between the demand and supply.

However, I will highlight to you that AI experience across from data center to edge, including in these edge devices like smartphones and PCs and other devices, AI experience, really more memory is essential. So without sufficient memory that AI experience, the functionality, the capability does get impacted in these edge devices as well.

So I mean, the punchline here is that AI across the board from data center to edge is driving increase in content and increasing requirement for memory as the customers look ahead at their roadmaps. And that's why customers are, of course, working with us with respect to access to supply for their long-term multi-year plans here as well.

Q - Vivek Arya {BIO 6781604 <GO>}

Thank you.

Operator

This concludes our Q&A session and today's conference call. At this time, I'd like to end today's conference call. Thank you all for participating. You may now all disconnect.

This transcript may not be 100 percent accurate and may contain misspellings and other inaccuracies. This transcript is provided "as is", without express or implied warranties of any kind. Bloomberg retains all rights to this transcript and provides it solely for your personal, non-commercial use. Bloomberg, its suppliers and third-party agents shall have no liability for errors in this transcript or for lost profits, losses, or direct, indirect, incidental, consequential, special or punitive damages in connection with the furnishing, performance or use of such transcript. Neither the information nor any opinion expressed in this transcript constitutes a solicitation of the purchase or sale of securities or commodities. Any opinion expressed in the transcript does not necessarily reflect the views of Bloomberg LP. © COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.